# 2026-04-28

Several PMs discussed the scope of the feature:

At first, Julian Smoller proposed a more extensive scope in which AI would actually do more steps:

**Proposal 1**

- In the SEC Filings search screen, we'd create a new entry point outside of the Protege chat, e.g. above the search filters. This could be an input field where users type or a button that they click, which might open a modal for them to type in (for example). It could even be a toolbar, like we used to have.
- When the user types and submits an input – e.g. "find 10-Ks that report a material cybersecurity incident" – AI would:
- (1) populate the search filters, e.g. Form = 10-K, Keywords = cybersecurity incident
- (see considerations below regarding whether we want to take additional steps)
- (2) execute the search – as a "normal" search, producing a tab of search results (same as clicking "Search" button)
- (3) Then, AI would proceed to analyze the search results one by one, producing a per-document analysis below each one.
- (4) Finally, it would produce an overview at the top of the results grid, summarizing what it found on that page of results.
- When the user navigates to another page (e.g. the next page) or otherwise changes the documents on the page, AI would analyze the new documents and re-create the overview.
- Note: All of this behavior (other than the initial input field) is identical to the prior version of AI search that we had in Protege v1, and the latter steps (where AI analyzes the search results and produces an overview) are identical to the "Ask about each document" feature that we currently have in v2.
- Note: At no point in the proposed flow would be open or interact with the Protege chat component. We'd frame this as an entirely separate "in-app" feature, just as other software companies have in-app features that exist independently of their chat / assistant AI.

But, we shifted to proposal two, which is now reflected in the "Requirements" section.

For the sake of preserving comments (which are otherwise lost when moving content to a new doc), we copied some of them here:

| |
|---|
| Jessica Powell - regarding the proposal to run a search<br>A data concern I had about the prior implementation: The search Protege did showed up as SearchLink and Search (the same search twice) in the history table, whereas now it should just show up as SearchLink. I could not easily see just user-generated searches in the prior implementation, which I would do to, for example, look at how users are constructing conceptual search or which filters are being used by users. Now I can just look for Search, I believe. This is mitigated somewhat for me if the user affirmatively selects search as I am more comfortable with that as user-generated, though there is certainly some room for debate there.  This is related to this debate but shouldn't be the deciding factor at all - just a data consideration re how we log in history for usage analysis. |
| Jessica Powell - regarding Input Triage and the other Questions & Decisions<br>I agree with the rest of your proposals - no triage, use a modal, don't worry about preserving the user input, SF home should still go to shelf, loading if needed for speed. |

Kristjan Bjornson notes:

Nice to have items:

- Insert AI's thinking somewhere
    - We have access to a rich data output, written in natural language, describing how to best use Intelligize filters to accomplish a search goal, based on expertise (e.g. "this info is typically found in.."). Note: We also have data regarding each filter individually, e.g. how AI wants to use a specific filter (but this is more granular and less important than the overall thinking / strategy).
- Suggest how to review the results (using a separate feature) to aid in the flow
    - How can we let our AI search tool know that there are other AI tools down the line that can help in answering a users question - eg 42 page example, we don't want our search AI to include 42 in keywords trying to narrow down documents that way. The optimal behavior would be to search for 10k's and then use another tool to analyze results for length
    - This is also important for our users to be aware of because AI may choose to do searches in a way that is confusing to our users

The more we think about the details of how this tool will run the less 'simple' it becomes. Should we move forward with an MVP product that only populates filters and does little else and then see how it works and how our users react?

~~~~

**Next Steps:**
~ Flesh out requirements in a separate doc (this one!)
(don't forget to fill out "project brief")
~ prepare a few data samples for designers to work with
~~~~

~ prepare a few wireframes to flesh out the ideas and jumpstart convo with designers
(and review the requirements as needed)
~ present to designers
~ (after a few design iterations)
~ present to engineers
~ finalize designs
~ estimate engineering time
~ schedule and begin dev work

# 2026-04-30

Overall, wireframes look good - KB to proceed with picking a design(s) to move forward with for the design team. In terms of the different ideas - probably a lower intervention approach to the UI is better eg don't mess with the filter drop down

Main note: from Julian, thinking steps (particularly thinking strategy) is high value to the user. Need to ensure that whatever approach we land on makes this visible/accessible to user.

Julian Smoller :
For example:
a button in filter panel, opens a modal
when modal closes, the button is replaced by "thinking" with loading icon and stop button
(when done, button comes back)
while thinking, an area opens up with the strategy (search advice)
it populates dynamically with text, using typewriter effect
user can close / re-open it
maybe it closes automatically when done populating filters?

…
Post Meeting Feedback:

- Difference between this and chat?
- UI clutter with three different entry points?
- Is there a good way to transition between this and chat?
- Why the extra steps? Why not search and run and analyze all in one?

2026-05-01

Julian: Just sharing a quick thought with @Kristjan Bjornson

(no need to respond - just brainstorming)
I was thinking more about the desire to have an end-to-end search & analysis flow in the app (what we did in v1). As discussed, the "help me search" tool might just fill out the filters, while the "ask about each document" tool analyzes the results. So, they'd both be granular and scoped to their part of the screen (filters, pre-search and analysis post-search). While I agree with this sentiment, I do still feel like it would be nice to have something that could do both together. Meanwhile, I was also thinking about the question regarding the search widget in the middle of the screen (in SF home). Maybe a compromise option would be to use that widget to do the full search and analyze flow in the app (as we did for v1). So, it would basically use the "Help me search" (builder) feature and the "analyze results" feature together in-app, instead of opening up the chat.

In fact, if we are really trying to rush something as a "fix" to the v2 feedback / desire for v1, we could implement that right now -- where we connect the home screen widget to the search and analyze flow (even before building the "help me search" button). As discussed, the team has all the code and logic already.

Just brainstorming!

2026-05-12

Meeting with design team following initial handoff of requirements:

Design mockup that shows modal/flyout - background is not greyed out. What will behavior be when user clicks outside the boundaries of the opened flyout?

- Before the user has hit submit: It functions like all other flyouts - clicking outside the flyout is the same result as hitting ‘x’ button. Flyout closes and resets.
- While ‘thinking’ is occurring, we need to understand what the behavior will be for the following three actions:
    - Click ‘x’ button
    - Cancel button
    - Clicking on the screen outside borders of the flyout

Of the two modal button options discussed - we like the one that is physically located in the filter panel rather than adjacent to search button, as it fits in more cleanly with user journey of filling out filters first then hitting search, when they go to fill out filters, we catch them and they use AI to fill out the filters for them

2026-05-15

Search Clarity with AI:

**Our stance today**: AI chat island needs to be separate from the intelligize UI. This is because if we want to do the things we want to do with AI (in the future) then it will not be compatible with intelligize UI (eg performing multiple searches across multiple apps and then comparing the results and returning only a subset). This is particularly true with the new architecture.

Important caveat: Technically under our existing architecture, things are still pretty compatible with the IGZ UI experience. And it wouldn't be very hard to integrate them actually (especially just the search part, analyze is more complicated).

Given that it would be very easy to integrate the results grid with V2 protege search, why don't we do it? People are asking for it.

- Philosophical opposition
- 'Training' our users on what a chat experience is or should be
- Creating space for embedded AI features now, so that we can build and introduce them, because eventually we really will need to divorce AI chat from IGZ UI
- Embedded AI features empower us to move more towards a world where AI exists in an island. IF AI did exist in a true island (fully disentangled and separate UI/Entry point) it would make it more clear the distinction between the different capabilities.

2026-05-20

Notes from discussion on miro wire frame between Julian Smoller and Kristjan Bjornson :

Julian (in black) Kris (in red):

Notes on Search Builder

Design:

- strive for consistency with other in-app action buttons (re: design, grammar, etc.)
    - common pattern: Verb Noun
        - Addressed. We will rename the button "Build Search"
    - common design: purple with rounded sides
        - Addressed
    - common icon: unique per action (status quo) or P for all AI actions (if we prefer, but then we need to change the other buttons to match)
        - Addressed. We will use a search icon instead of protege P. See Miro mockup for details. Link above.
    - common placement: to the right of the "normal" (non-AI actions)
        - Addressed
    - the modal itself should also have the same name, e.g. "Build Search" (instead of Search Builder)
        - Addressed

- thinking state
    - ideally we should avoid using the button itself to represent the loading state (and stop button)
        - Addressed - we will have a loading state similar to other loading states in intelligize for both the search panel and results grid
    - there are a lot of other actions that take time to complete, and we have existing design patterns that we should use, instead of adding a new one
        - Addressed see above
    - e.g. in the modal itself
    - e.g. the filter area is in a loading state
    - I recommend asking around / exploring our product, and other products, to find relevant patterns - how we tackle similar situation in other contexts
- dynamic button - Thinking
    - I agree that we need a component like this in this general location.
    - Consider that, structurally, this is NOT like an action button. It is something that shows / hides additional content, like the info icons we show next to each filter today. We should explore additional (existing) design patterns for this type of component.
        - Addressed - but not sure about the solution. See miro mockup design. Tried to implement something similar to overview functionality in ‘ask about each document’
- dynamic buttons - Ask
    - We should NOT add a new, dynamic "Ask" button in the filter panel, because that would break a bunch of structural rules. Instead, we should use the existing Ask button (on the right) or consider other options
    - Options:
    - Use the pre-existing ask button, on the right, and:
        - pre-populate the instructions?
            - maybe clarify that these are based on the user's search
        - draw user's attention to it via quick tip?
            - Addressed. New design includes check box within modal to skip straight to analysis. Also added a quick tip design element for ask about each document to prompt the user to complete analysis.
    - Within the initial modal, include a checkbox or additional button to trigger analysis aftter search, e.g.
        - "Automatically analyze search results" (checkbox)
            - Addressed see above

Structure:

- consider when exactly we open a new tab and the implications of that
    - When user submits search request to aI, e.g. clicks "OK" in the modal?
        - This one makes the most sense for reasons explained below. It fits into our existing design patterns the best, and gives us the ability to avoid creating stop button which would be tricky. We will pursue this avenue for now until we find a reason not to.
    - or after AI populates filters and executes search (as is usually the case when user themselves clicks on Search)
    - Consider what exactly screen would look like in the intermediate states
    - Proposal:
    - - immediately open a new tab when user submits request .. in which case we don't have to worry as much about a stop button; user can just close the tab instead.
    - filter panel and results grid might be in loading state -- could be weird? but the alternative could be even weirder, e.g. showing prior search results
    - (could even use that area to show thinking, before it collapses into a thinking area -- although that would break the structural rules)

# 2026-05-29 - Engineering

Thoughts on engineering work:

Rely on current search flow? No…

- We should just create a separate prompt and flow for this feature, instead of relying on the exact same steps we currently have for the conversational search. Why?
- We won't have a manager step, so the input to the strategy step needs to be a little different; instead of a manager-generated search instruction, it will be what the user wrote, and so we need to add some guidelines around how to interpret it.
- We need to produce slightly different outputs for different purposes and priorities. Specifically, we primarily need to produce one user facing strategy output, instead of the multiple outputs produced by the current step.
- We want to prioritize, and hopefully we can find a way to speed things up (maybe a different model, or maybe just simplified the prompt and outputs).
- Longer term, the current search flow (for conversational AI) will probably go away, actually, because we'll be shifting to new architecture, so we don't want to rely on it as a long term thing.

Steps we want to run:

- (skip manager / input triage)
- Search Strategy - generate strategy & analyze instructions
    - May need to adjust, given need to interpret user inputs

    - Need to adjust to produce different outputs & prioritize speed
    - Can we load this into the UI before anything else?
- Search Planning - select filters to use and generate filter instructions
    - (maybe this could be combined into prior step?)
- Search Building - select filter values and compile into executable search (link)
    - (separate step for each filter that is used)
    - Load filters into the UI
- Run Search - (non AI step)
    - Load search results into the UI
- (if the checkbox is checked to automatically analyze:)
- Analyze each result - analyze each result independently to produce a per-document analysis
    - Load per-document analyses into UI
- Generate overview - consolidate analyses of each result into overall overview
    - Load overview into UI

Speed:
- As part of this project, we should take a closer look at which steps take the most time and how we can optimize them for speed.
- Merge planning with strategy?
- Break out analyze instructions from strategy?
- Different models?
- …

"Search Comments"
- FYI we implemented something to deal with situations where AI:
- a) decides to use a filter that has not been integrated with AI yet
    - This is rare, and perhaps was a bad idea in the first place, but we basically give AI a complete list of the filters in an app, even though some of the lesser used ones are not supported, so it's possible that AI decides to use a filter that it cannot use. In that case, it generates a "comment".
- b) encounters an error when using a filter, e.g. it can't find the value it wanted to find
    - This is less rare. Sometimes AI will plan to use a filter because it expects to find a certain value, e.g. a company within the company filter, or a company type within the company type filter, etc. But when we get to the "search building" step and give AI more information about the filter, including a set of values in some cases, it may realize that it can't do what it initially planned to do, and so instead of blocking the search from proceeding, we still create and run it, but we generate an additional "comment".
- All comments produced during the AI steps are aggregated into an overall user-facing "search comment" that is displayed as part of thinking steps in v2 (when it exists).
- Is there going to be a similar concept for this project?

    - Option a) Don't generate or show a search comment anywhere, but still produce and run a search (which will be deficient in some way), and treat this as a data quality issue to prioritize and improve relative to others
    - Option b) Generate a search comment, like we currently do, and find a home for it in the UI. (But note: even with the comment, the search is still deficient.)
    - Option c) Change the strategy and planning steps to avoid telling AI about filters it cannot use, which will prevent at least one of these scenarios from happening.

Automatically Analyze:

- The current plan is to use a checkbox to allow the user whether to launch into the analysis.
- Perhaps we should always generate the analyze instructions in the prompt, though, but only apply them if the checkbox is checked? (As opposed to changing the prompt.)
- Even if we don't auto-analyze, we may want to use those instructions to populate the "review results" field.
- But does it slow down the AI processing? Maybe.
- Could we split the generation of analyze instructions into a separate step? Maybe…
- Is it possible that AI might change the search that it does based on whether or not it plans to analyze the results?

Consider patterns:

- Primary Pattern - The search conditions are covered by filters, including keywords, and can be confirmed with post-search analysis.
    - Find 10-Ks that disclose a material cybersecurity incident >> Form: 10-K, Keywords: cybersecurity incident, Analyze Instructions: confirm that this 10-K discloses a material cybersecurity incident (because keywords don't guarantee it)
- Conditions not covered by filters - At least some of the conditions in the search request cannot be covered by any available filters.
    - Find 10-Ks that are over 42 pages long
- Find a doc and extract info from it ….

# 2026-06-03

Tentatively, we want to:
- create a data design – i.e. a series of steps (powered by LLMs) that make decisions and produce what we want
- for each step, create / edit the relevant prompt
- testing & refinement

2026-06-04

KB to start project docs for other embedded AI features

Having documentation along with designs to present to internals on create search, justify design decisions/layout

# 2026-06-12

Create search kickoff notes:

FIrst goal - who will work on create search from NBS team?

-Muhib - UI
-Saim - backend changes
-Uzair - AI

KB explains scope of project to team

Now let's estimate timing of this project

-NBS owes an estimate of timing on this project

# 2026-06-16

By Ehtisham:

After reviewing the requirements, design details, and Figma screens for this project, there are a few areas that appear ambiguous and may require further clarification and discussion.

At a high level, from a development perspective, we may be able to leverage parts of the existing **"Run Search"** implementation. However, it is unlikely that we can reuse the flow as-is due to differences in the expected behavior and user experience.

For example:

- In the existing **"Run Search"** flow, the search filters are displayed within the chat response, and the user explicitly clicks a button to execute the search. In the new flow, we need to display a different type of chat response, automatically populate the filters, and trigger the search without requiring any user interaction.

- The filter structure received from the backend typically requires parsing and transformation into the format expected by the frontend components before the filters can be populated. Similar handling may be required for the new AI-driven search flow.
- There may be additional differences in the response structure, state management, and execution flow that need to be considered.

Based on these differences, we may need to introduce a dedicated flow for the **AI Create Search** functionality while reusing existing components where applicable. This approach would allow us to minimize duplication while avoiding unnecessary complexity or potential impact on the existing search implementation.

# 2026-06-16

By Uzair:

As per our current understanding of the requirements for this project, there are a few areas which are a bit ambiguous & might require more discussions, even POCs & experimentation:

1. In Protege-V2, we generate <analyze-instructions> from the Search Strategy step using only the user's query as input. In Protege-V1, we generated a review goal after search construction, using both the user query and the generated search. Given that Create Search will support both suggested analysis and auto-analysis, should analysis instructions be generated directly from the user query (V2-style), or should we introduce a separate V1-style review goal step that has access to the constructed search and selected filters?
2. Since speed appears to be a primary goal for this feature, it may be worth exploring whether the Search Strategy and Search Planning steps can be combined into a single prompt. This could potentially reduce latency by eliminating an LLM call while still producing the required outputs (research strategy, search instructions, planning outputs, and potentially analyze instructions). We would likely need some experimentation to understand the impact on search quality and prompt complexity.
3. Regardless of the final workflow, we may want to evaluate faster models (e.g., Gemini Flash, GPT Instant, Haiku) for some or all LLM prompts/steps, as the requirements place a strong emphasis on speed, responsiveness and progressive loading of results.

Based on the above points there can be 3 flows:

## Flow 1 - Closest to V2 Flow

**Steps:**

**User Query → Strategy → Planning → Building → Search Execution → Ask about each doc → Consolidation**

- **User Query:** Raw user input.
- **Search Strategy** takes the User Query as input and generates:
    - `<research-strategy>`
    - `<search-instructions>`
    - `<analyze-instructions>`
    - `<user-message>`
- **Search Planning** takes the generated `<research-strategy>` and `<search-instructions>` as input and generates a plan containing all suggested filters and filter instructions.
- **Search Building** takes the suggested filters and filter instructions as input and builds the search by populating actual filter values.
- **Search Execution** executes the generated search and loads results into the UI.
- **Ask about each doc** step (if applicable / selected by the user) takes the generated `<analyze-instructions>` as input and generates per-document analyses.
- **Consolidation** generates a consolidated overview based on the User Query and the results of the Ask about each doc step.

**Pros:**

- Closest to existing V2 implementation.
- Lowest engineering effort and implementation risk.
- Existing prompts and orchestration can largely be reused.
- Already proven in production.
- Easier to maintain and debug due to clear separation of responsibilities between steps.
- Analyze instructions are always available.

**Cons / Considerations:**

- Highest latency due to multiple sequential LLM calls.
- Analyze instructions are generated before search construction and therefore may have less context than a search-aware review goal.

---

## Flow 2 - Merged Strategy + Planning in the V2 Flow

**Steps:**

**User Query → Strategy + Planning → Building → Search Execution → Ask about each doc → Consolidation**

- **User Query:** Raw user input.

- **Strategy + Planning** takes the User Query as input and generates:
    - `<search-instructions>`
    - `<analyze-instructions>`
    - `<user-message>`
    - Suggested filters and filter instructions
- **Search Building** takes the suggested filters and filter instructions as input and builds the search by populating actual filter values.
- **Search Execution** executes the generated search and loads results into the UI.
- **Ask about each doc** step (if applicable / selected by the user) takes the generated `<analyze-instructions>` as input and generates per-document analyses.
- **Consolidation** generates a consolidated overview based on the User Query and the results of the Ask about each doc step.

**Pros:**

- Removes one LLM step compared to Flow 1.
- Faster overall search generation.
- Simpler orchestration and workflow.
- Reduces end-to-end latency while preserving most existing functionality.
- Analyze instructions are available immediately after the first AI step.
- Potentially a good balance between implementation effort and performance.

**Cons / Considerations:**

- Requires creation of a new combined Strategy + Planning prompt.
- Requires experimentation and prompt tuning to maintain search quality.
- Less modular than the current V2 architecture.
- Harder to independently optimize Strategy and Planning behavior in the future.
- Analyze instructions are still generated before search construction and may lack search-specific context.

---

## Flow 3 - Merged Strategy + Planning with Separate Review Goal (from V1 Flow)

**Steps:**

**User Query → Strategy + Planning → Building → Search Execution → Review Goal → Ask about each doc → Consolidation**

- **User Query:** Raw user input.
- **Strategy + Planning** takes the User Query as input and generates:
    - `<search-instructions>`

    - `<user-message>`
    - Suggested filters and filter instructions
- **Search Building** takes the suggested filters and filter instructions as input and builds the search by populating actual filter values.
- **Search Execution** executes the generated search and loads results into the UI.
- **Review Goal** takes the generated search information as input and generates:
    - `<analyze-instructions>`
- Inputs:
    - Search filter values
    - Search instructions
- Output:
    - Analyze instructions
- **Ask about each doc** step (if applicable / selected by the user) takes the generated `<analyze-instructions>` as input and generates per-document analyses.
- **Consolidation** generates a consolidated overview based on the User Query and the results of the Ask about each doc step.

**Pros:**

- Faster search generation than Flow 1 due to merged Strategy and Planning.
- Review Goal can be generated in parallel to Search Execution.
- Review Goal can generate analysis instructions using actual search context.
- Produces more search-aware and potentially higher-quality analysis prompts.
- Closer to the V1 approach for individual document analysis.
- Better alignment between generated search and document analysis.

**Cons / Considerations:**

- Requires creation of a new combined Strategy + Planning prompt.
- Requires implementation of an additional Review Goal step.
- More implementation effort than Flows 1 and 2.
- Additional orchestration complexity.
- May require experimentation to determine the best inputs for Review Goal generation.

# 2026-06-29

Initially, we will have a smaller more focused team working on create search

Backend: Saim
UI: Muhib
AI/Data: Shawana

KB review of notes from Ehtisham and Uzair:

Ehtisham:

Points acknowledged, let's create a distinct flow, starting with as many components from the existing run search flow as possible. We certainly want to build something that operates independently of the existing run search flow. We will likely want to start with the building blocks we already have, asses performance and then make changes as applicable.

Uzair:

I appreciate all these points.

Let's discuss pros/cons of generating analysis before or after search. Points of clarification - is analysis goal in V1 generated based on the search constructed only? Or the search constructed as well as the results of said search? Why did we change this for v2?

In terms of the three flows you discussed - do they all have a similar or the same starting point? From what I can tell, It seems that at least for flows 1 and 2, we could start with the simplest and fastest option from an engineering perspective (flow 1) and then iterate from there towards flow 2. The main difference in flow 3 is when the analyze goal is generated. My thought here is that we should start with flow 1 and then iterate towards flows 2 and 3 depending on performance of flow 1. Let's discuss if this approach makes sense.

For tomorrow: 6/30 - Let's start with trying to agree on some basic organizational principles for this project.

Let's create a narrative (or could be a flowchart) that outlines the steps that occur in our create search flow. For each call to LLM API that includes a prompt - we have a link to a relevant prompting document.

I am looking at this [doc ]for inspiration

# 2026-06-30

Open question: Is it better to generate analyze instructions as part of the search strategy (or in parallel with) or after we have generated the search terms?

Search comments - for example, will we provide search comments like we do today eg surrounding peer lists etc. A larger question is will we provide any sort of commentary or user feedback in create search? How would we do this given that create search is non-conversational?

**Meeting summary**: We discussed two main items today - Picking a flow (from options above), and documentation

Picking a flow: As discussed during our call today, there are a variety of options in terms of what flow we use and how we implement. Likely any of the options discussed are acceptable so the engineering team can use their discretion in terms of picking and beginning to implement a flow. We want to move quickly on this project, so ideally, we will reuse components from protege where applicable.

Documentation: As we pick a flow - let's create a narrative (or flowchart) that outlines the steps that occur in the flow, and the prompts we are using at each step. Please follow the existing prompt documentation format when documenting prompts (e.g. each prompt should have its own prompt doc with documentation, goals, etc.). Ideally, this narrative should be a living document that explains how our create search flow works at any given time.

For next time: Engineering will pick and propose a flow - we can discuss any questions or concerns associated with the flow and begin the documentation process.

# 2026-07-01

## Feedback

Meeting with nicole:

Will they see their initial prompt after submitting

Can they iterate on create search using AI? Eg could a user run a search with create, and then open the modal a second time to iterate on their initial search?

Desire from Nicole is to either populate original prompt somewhere, so that a user can iterate using AI/Create

Nicole asks for the ability to save the default number of results displayed on a page as a setting. This would impact the functionality of create search, especially in the context of have the analyze search results box checked.

What about the name and logo of the button?

> Nicole comments: can we get rid of the clear button?

If a user decides to check the analyze box, how can we get them to input not just what they are searching for, but their analysis goal as well? We may need to update the copy of the description for the modal, or even create some kind of dynamic instruction

Search strategy should be collapsed by default. Maybe make it easier to expand. Just clicking the box should expand it.

Nicole's main point - desire for ability to iterate

# 2026-07-02

Meeting notes:

Team has decided to implement flow #1 based on ease of implementation and the desire to move quickly

**Strategy Step:**
Of the steps in v2, the strategy step is likely the step that will require the most attention from us. This is because the strategy step in v2 receives information from the manager step, whereas here, the strategy step will receive a raw input from the user. Also the outputs may change - today we generate research strategy and user message, as outputs (in addition to search and analyze instructions), we may want to update or modify these outputs to be better suited to this specific use case. It's also possible we may want to update analyze instructions as well to be more specific to this use case.

- Inputs are changing
- Outputs are changing

Search Strategy Prompt:
https://docs.google.com/document/d/1IF5Z-hpLO4WpYL76iDlb21tmu2hGREA48jlX5VrQ4PI/edit?tab=t.0

**Planning step**: This step is focused on creating a specific output, like what specific filters to use and what to search for within each filter

**Search building:** We currently don't plan to change anything in these steps, so we will ignore them from our discussion for now.

**Next steps:** For Monday - let's have a focused session on strategy prompt.

KB to come up with initial test set to begin evaluating prompt quality

**Open question discussion**:

Analyze instruction generation location:

Having analyze later could improve speed to search execution (it moves the step to after search execution). Conversely it could slow down the execution of the entire flow (though that might not matter much if we stream results).

More important - the generation of the search strategy may be dependent on the ability to analyze later (eg searching for a specific page number within a 10k). Therefore, giving AI the ability to analyze at the same step where it is generating a search may improve search quality.

One thing to note is that the relationship between search construction and analyze is fraught today, Harry observes lots of issues with this eg protege says it is going to some analysis

Also, important to note that

Regina feedback meeting:

The button may need provenance or additional attention. Inexperienced users may miss the placement of the button. Marketing and training is going to have to really push the location of this new button. Button isn't as obvious as it could be.

Okay with the name

Analyze search results checkbox is intuitive

Regina thinks the link between ask and create are relatively intuitive?

CB feedback - interactive concern, maybe we just load previous prompt when create button opens after initial search

Reduce height of search strategy overview in both expanded view (padding on bottom), and minimized view (could be - Search strategy: Iran war risk factors). Height/space on results grid is very important to users

Need to make popup for ask about each doc bigger

# 2026-07-06

Added By Shawana Faz

**Prompts Details:**
Prompts for create a search are added in this folder: Create a Search Prompts.
Here are the prompts and the flow of prompts that will be used in Create a search.

For Flow 1: **Prompt Search Strategy** → **Prompt Search Planning** → **Prompts Search Building** → <<Step Search Execution>> → **Analyze Instructions/Review Goal**

**Dataset**:
We prepared a dataset for testing the prompts for create a search.
2026-07 - Create a Search

**Queries:**

1.Will we require the search strategy/planning generated in this case to be logged somewhere? Or would we require only the search build to be stored?

2.For thinking steps, would it require each step, like strategy,planning to be shown on UI? Or only visibility of strategy as in the figma designs?

3.If once analyze instructions are executed, will the ask about each document modal still show the ran analyze instructions. In case of an auto analysis run, where will we show it.

4. Will there be any triage/error handling for out of scope queries?

# 2026-07-06

High level Flow Diagram based on the **Flow 1 - Closest to V2 Flow.** Added by Saim Zahir

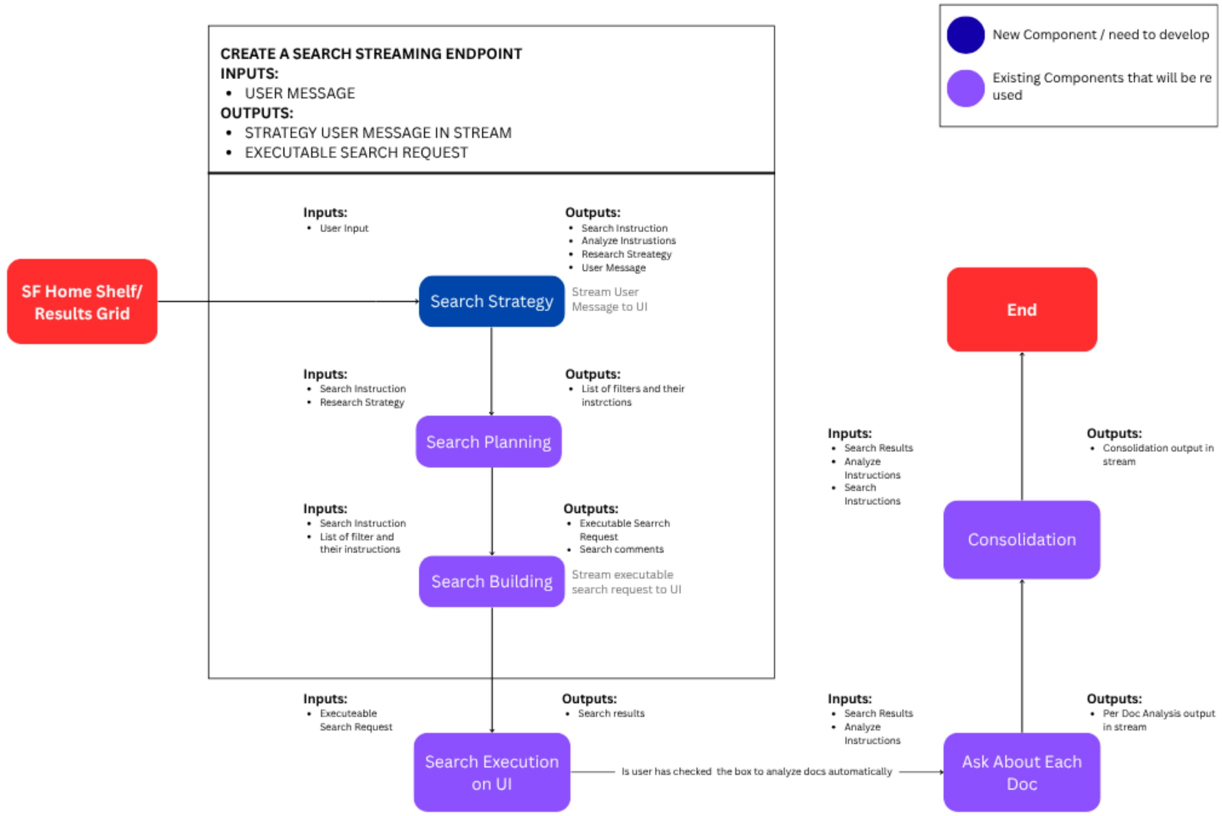


Note: There could be no results found in Search Execution step - consider what happens then (e.g. where do we show strategy?)

Consider some things to be adjusted for in-app flow (v. chat):

**Search comments:**
Plan A: Do nothing
Plan B:
Maybe we show a comment / written commentary on the search, as part of / after the strategy?

**Input triage:**
Consider:
1) vague requests, like “x”
2) clear but out of scope requests, like “find SEDAR filings”

Plan A: Run a search as best we can (no triage, user inputs -> strategy)

Plan B: “light” input triage that runs while modal is still open, yielding a rejection in some cases
Shawana mentions that we can get rid of the manager prompt but still retain some input triage step - could use a fast model and still be lightweight.

Main challenges with input triage:
Items that are nonsensical eg ‘x’ or gibberish
Items that make sense but are clearly out of scope - eg look for SEDAR filings

**Analyze Instructions → Ask about each document**
Note that there could be a potential mismatch between the Analyze instructions (output by the strategy step) and the expected input for the “Ask about each document” flow.

**To Do:**
>> Make prompt files for whatever steps run within “ask about each document” flow

**Open Questions:**
- Consider combining first two steps
- Do we need to generate two separate research outputs in step 1? Or can we just generate one that serves two purposes?
- Should analyze instructions be generated in step 1 or separately?
-

**Decisions**:
- We will keep the analysis generation within the search strategy step, we will add context to the prompt to make sure that AI knows that analysis is optional and won’t always occur
-

Julian: Need to drop but quick notes to your Qs:

1) we mainly need the logging to allow us to troubleshoot what went wrong - possibly could do this via testing outside of production, but would be helpful to have logging for production as well. (Could just be s3 or something, if not SQL database)

2) no thinking steps required, outside of strategy

3) not sure

4) TBD - up to KB - as discussed, could be a light version of input triage that occurs in modal

# 2026-07-07

**Agenda:**
-”Ask about each document” prompts and flow overview

Added by Saim Zahir

## Overview of Ask About Each Document

The Ask About Each Document feature generates an individual response for each selected document based on the user's Input.
**Input:**
- User Input
- Selected Documents from Search Results

**Output:**
- For each selected document, the process generates an independent AI response based on that document's content

**Content Retrieval**
- Check whether semantic chunks are available for the document.
  - If semantic chunks exist:
  - Retrieve up to 300 candidate semantic chunks.
  - Apply a re-ranking step to identify the 15 most relevant chunks for the user's question.
  - Use these selected chunks as the input to the AI model.
- If semantic chunks are not available:
  - Use the entire document content as the input.

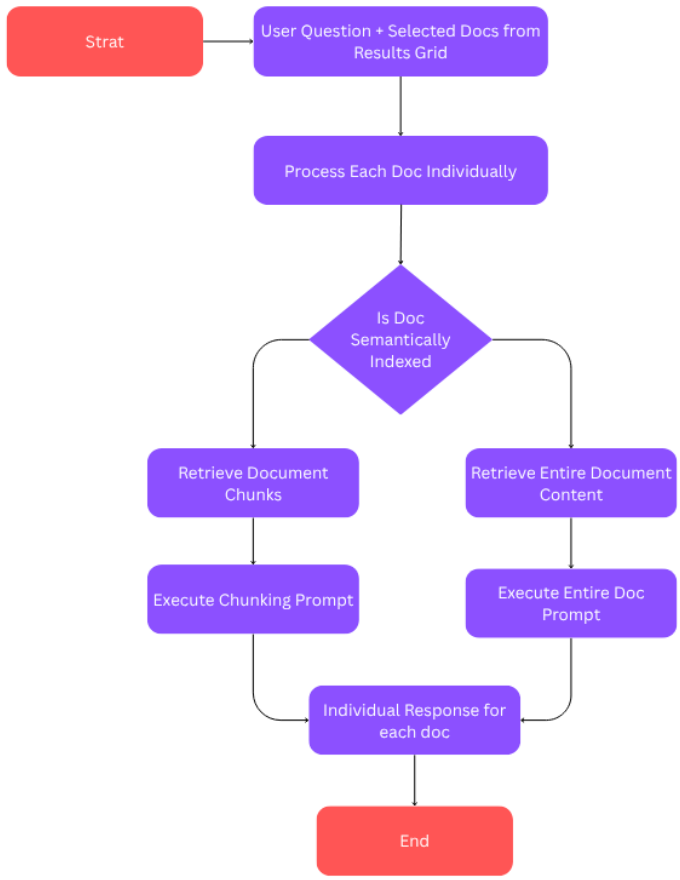


## Overview of Consolidation Process

The **Consolidation Process** runs after the **Ask About Each Document** feature completes its analysis. It combines the individual document responses into a single, consolidated result based on the user's original query, providing a unified view of the findings across all analyzed documents.

**Inputs:**

- User's Original Input
- List of analyzed docs including
  - Company Name
  - Form Type
  - Document Id
  - Analyze result

**Output:**

A single consolidated response that summarizes and synthesizes the findings from all analyzed documents in the context of the user's query.

2026-07-09

**Agenda:**
-"Ask about each document" prompts and flow overview
-Strategy Prompt discussion 2026-07 - AI Prompt - Create a Search - Search Strategy
-Updated planning prompt as well 2026-07 - AI Prompt - Create a Search - Search Planning

KB:

Review of ask about each document flow:
-Need to consider whether we want to use chunking or whole document review for ask about each document
-Julian added backlog item
-Also, in apps where we don't have v2, ask about each document is just 'ask AI' which includes a triage step.

One of the major updates that needs to be made in the search strategy prompt is the generation of the analyze instructions. Consider the following:

- As discussed above, we may update the downstream flow (ask about each document), in two different ways
    - By removing semantic chunking from the ask about each doc flow
    - By updating the whole document prompt to be more in line with our current analyze prompt(s)
- Ultimately, these changes shouldn't or won't affect how we structure our prompt for create search - we should likely just ignore the semantic chunking and try to generate good/workable instructions for whole document flow
- One key consideration of the downstream prompt is whether the analyze instructions that get generated should be stated as per document instructions, or whether they should be stated as a generic goal
    - Eg - "for each document, perform the following analysis steps" vs. "What are the main risk factors I should consider disclosing based on my peer companies?"

# 2026-07-13

**Status Update By** Saim Zahir
The basic end-to-end implementation for Flow 1 has been completed on the backend. Now discussing with the UI person to integrate the endpoint for further testing.

- Created a separate endpoint to accept **user input** and return the **Search Request(executable search)**

**Status Update by** Shawana Faz
For Flow 1, we've ran an iteration with the discussed changes to strategy prompt, and are getting good outputs. Therefore no further prompt changes to be made.
Once integrated with UI, we'll test the dataset end to end, and check for any issues/tweaks to be done to prompt. 2026-07 - Create a Search

**Status Update by** Muhib ur Rahman Bakar
- Added "Create Search" button + popup for entering a search query in plain language
- Wired to a mock API (demo/placeholder only) that simulates an AI-generated search strategy and filters
- Filters panel shows a loading skeleton while filters are being populated
- Results toolbar appears instantly (disabled) with a collapsible "Search Strategy" panel showing loading progress

KB discussion points:
- Speed?
- User message in parallel
- Output format of prompts
- Evaluation of analyze instructions
- Testing, what have we done, what more can we do

Meeting notes:
- When a user saves or restores a search from create search, should we preserve search strategy and/or analyze instructions?
    - For now, no. This would open up a lot of potential complexity. KB will note this down as a question/decision and we will move on. We can explore this as a possibility going forward but not currently considered part of MVP.
- Can we remove spacing around search strategy - to free up more space? Potentially include in upper section like the toolbar directly above it
- Preserving prompting - do we need to preserve somehow, the user input so that a user can iterate? Or even more complex, should create search be somehow aware of previous inputs/searches so that a user can iterate?
-

# 2026-07-20

Topics:

~ (see v3 doc for other stuff, e.g. visibility, AOE)
~ create search status

**~~ Create Search**

We are making prompt changes, but also:
Speed - we are looking into it

Prompts
- strategy prompt → many changes, ….

~~~~~~~~~

Ideally, the prompt-writing workflow goes like this”
0) list the prompt in the prompts sheet and create a dedicated doc using template
1) first, in a documentation tab, write up all the relevant context about the prompt that we can think of
2) then, give the initial documentation to AI (via GAI chat), and ask it to respond with questions, suggestions, push back, feedback, etc. – in other words, to talk about it before jumping into writing anything
3) assuming AI responds with interesting questions and suggestions, review them, decide on certain things, and write down the decisions / thinking / responses to AI questions (in the documentation tab, and then in the chat itself)
4) then, ask AI to prepare clean documentation based on the initial documentation we sent it and then subsequent conversation with questions / answers / feedback / etc
5) continue iterating on documentation as needed - for example, if we realize something important is omitted, we can first explain in our own words, then ask AI to update documentation, etc.
→ **>> clean and clear AI written documentation about the prompt / LLM step**
the “latest and greatest” AI written documentation should be in a documentation tab, ideally with a link to the chat in which the documentation was generated (hopefully with a high powered GAI model)

6) then, in same chat (or possibly a new chat where we paste the documentation), we ask AI to generate an initial draft of the prompt
- Note: We are NOT giving AI our current draft of the prompt. Why not? Because we want to see how it drafts the prompt based purely on the documentation and requirements.
7) inspect the AI drafted prompt and compare it to any versions that we drafted / are using. For any differences that we observe, consider if:
a) maybe the AI version is better
b) if we like our version better, then consider if maybe AI is omitting something because of a gap in the documentation, in which case we can circle back and update the documentation
~~~~~~~~~

8) usually, AI will comment on its draft of the prompt, e.g. suggesting possible enhancements, or at least pointing its reasoning for certain things. This commentary can be very valuable and lead to additional learnings and/or improvements on our end, so it is good to read and react and document it.
~~ we may respond to an AI suggestion by saying, “Yes, go for it,”in which case it might edit the prompt a bit more, so there could be some iterations here, but either way, we want to get to:
**→ >> clean and clear AI written version of the prompt**
– this should be presented in a proposed prompt tab in our prompt doc, along with any AI commentary on the prompt, and ideally with a link to the chat used to generate it

**To Do:**
- run this process for search strategy prompt (used in Create Search)
>> get to AI written documentation of strategy prompt
>> get to AI draft of this prompt

# 2026-07-27

Status??

Data quality/testing

## Speed

An important topic and core requirement

- Where are we at now? One question is what is the relationship between the demo internal environment and production environment?
- Let’s assume that our demo internal environment is close enough to production - so we’ll start there
- First, we need to decide if speed is even something we need to improve - so we need an accurate measure of latency over our test set.
-

Added by Shawana Faz
2025-12-24 - Latency Testing Search Flow V2 These are the stats we prepared for V2 launch on demo internal. We investigated which step took time and then further added detailed logs on it. And did improvements.